

Kubernetes Storage Architecture: Volumes, PersistentVolumeClaims (PVCs), and VolumeClaimTemplates (VCs) in kubectl

1. Overview

Modern cloud-native applications often require persistent storage that survives container and pod restarts. Kubernetes, being a container orchestration platform, provides a robust storage framework to address such needs. This document explains the foundational concepts of **Volumes**, **PersistentVolumes (PVs)**, **PersistentVolumeClaims (PVCs)**, and **VolumeClaimTemplates (VCs)**, specifically in clusters set up using **kubectl**.

2. Kubernetes Volumes

2.1 Definition

A **Volume** in Kubernetes is a directory, possibly backed by storage media, which is accessible to the containers in a pod. Kubernetes volumes extend the lifecycle of storage beyond container restarts within a pod, unlike Docker volumes which are tied to a single container instance.

2.2 Characteristics

- Volumes are mounted into containers at a specified path.
- A volume's lifecycle is bound to the pod; it exists as long as the pod exists.
- Data is preserved across container crashes but not across pod deletions (except with PVs).

2.3 Common Volume Types

Type	Description
<i>emptyDir</i>	Temporary storage shared across containers in a pod. Deleted with the pod.
<i>hostPath</i>	Mounts a path from the host node. Limited to development/test environments.
<i>configMap</i>	Injects configuration files into a pod.
<i>secret</i>	Provides sensitive information like passwords or tokens.
<i>nfs</i>	Mounts a network file system into the pod.
<i>csi</i>	Allows integration with external storage drivers via the CSI specification.

3. PersistentVolumes (PV)

3.1 Definition

A **PersistentVolume** is a cluster-scoped resource that represents a piece of storage provisioned by an administrator or dynamically by Kubernetes through a **StorageClass**. It is independent of any individual pod and persists across pod lifecycles.

3.2 Features

- Pre-provisioned by an admin (static) or created on-demand (dynamic).
- Defined using a YAML configuration with capacity, access modes, and storage source.
- Managed by the control plane component kube-controller-manager.

3.3 Access Modes

Mode	Description
------	-------------

ReadWriteOnce	Mounted as read-write by a single node.
---------------	---

ReadOnlyMany	Mounted as read-only by many nodes.
--------------	-------------------------------------

ReadWriteMany	Mounted as read-write by many nodes. Supported by some CSI drivers.
---------------	---

4. PersistentVolumeClaims (PVC)

4.1 Definition

A **PersistentVolumeClaim** is a request for storage by a pod. It is a user-facing resource used to abstract away the details of how storage is provisioned and used.

4.2 Lifecycle

1. A PVC is created by a user/application.
2. The control plane searches for a matching PV that satisfies the claim.
3. Once bound, the PVC and PV remain bound for the life of the claim.

4.3 PVC YAML Example

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
metadata:
```

```
  name: example-pvc
```

```
spec:
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
  resources:
```

```
    requests:
```

storage: 1Gi

storageClassName: standard

5. VolumeClaimTemplates (VCs)

5.1 Purpose

VolumeClaimTemplates are used primarily with **StatefulSets**. Each replica of a StatefulSet gets a uniquely bound PVC created from the template. This ensures persistent, identity-aware storage across pod rescheduling.

5.2 Key Benefits

- Unique storage per pod.
- Automated PVC creation per replica.
- Essential for applications like databases, queues, and stateful services.

5.3 Example in StatefulSet

volumeClaimTemplates:

- metadata:

name: data

spec:

accessModes: ["ReadWriteOnce"]

resources:

requests:

storage: 5Gi

storageClassName: standard

6. Dynamic Provisioning and StorageClasses

6.1 What is Dynamic Provisioning?

Dynamic provisioning allows Kubernetes to automatically create storage volumes based on a PVC. This requires a **StorageClass** which defines how the storage should be provisioned (e.g., through AWS EBS, GCE PD, or a CSI driver).

6.2 Example StorageClass

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: standard

provisioner: kubernetes.io/aws-ebs

parameters:

type: gp2

6.3 Important Considerations

- kubeadm clusters do not come with a default StorageClass.
 - You must deploy a CSI driver (like hostpath CSI for local testing or AWS EBS CSI for production).
 - Persistent storage in production should always use a managed and replicated storage backend.
-

7. kubeadm Considerations

7.1 Default kubeadm Setup

- No dynamic provisioner is deployed by default.
- You may need to install a CSI driver manually.
- Ideal for bare-metal or custom environments where you want full control.

7.2 Recommended CSI Drivers

Environment	CSI Driver
Local testing	hostpath CSI driver
AWS	aws-ebs-csi-driver
GCP	gcp-pd-csi-driver
Azure	azure-disk-csi-driver
On-prem storage	Ceph, OpenEBS, Longhorn, NFS CSI

8. Best Practices

- Always use PVCs to request storage in pods, not PVs directly.
- Use StorageClass for dynamic provisioning instead of managing PVs manually.
- In production, ensure high availability and backup of persistent volumes.
- Avoid hostPath volumes in production due to node-dependence and security concerns.
- Use VolumeClaimTemplates with StatefulSets for applications requiring unique, stable storage.

9. Conclusion

Kubernetes provides a robust and extensible framework for managing storage through Volumes, PVs, PVCs, and VCs. When using **kubeadm**, administrators must ensure proper setup of storage provisioners and CSI drivers. Understanding these storage constructs is crucial for deploying and maintaining stateful applications effectively in a Kubernetes cluster.