

Kubernetes Theoretical Concepts

Pods

A Pod is the smallest and simplest unit in the Kubernetes object model. It represents a single instance of a running process in your cluster.

Key Points:

- A Pod can contain one or more containers (most commonly one).
- All containers in a Pod:
 - Share the same network namespace (IP address and port space).
 - Can communicate with each other using localhost.
 - Share storage volumes if defined.
- Pods are ephemeral. If a Pod dies, Kubernetes does not restart it automatically unless it's managed by a higher-level object like a Deployment.
- Pods are typically created and managed by controllers (like Deployments, ReplicaSets, StatefulSets).

Pod Lifecycle:

- Pending: The Pod has been accepted but not yet scheduled.
- Running: The Pod is bound to a node and containers are running.
- Succeeded: All containers have terminated successfully.
- Failed: One or more containers terminated with an error.
- Unknown: Pod state cannot be determined.

Commands:

- Use command to install kubectl
sudo apt install kubectl
- To create a pod, navigate to a specific namespace or activate namespace using:
kubectl --namespace="namespace_name"
- To run the pod use command *kubectl run*
Example:
kubectl run nginx --image=nginx //the pod will be created
- *kubectl get pods*
- *kubectl get pods -n nginx* //where "nginx" is a namespace
- To delete the pod from default or current namespace:
kubectl delete pod "pod_name"

- To delete the pod from a specific namespace:
kubectl delete pod "pod_name" -n "namespace"
 - To create pods make a *pod.yml* file
 - *kind: Pod*
apiVersion: v1
metadata:
name: nginx-pod
namespace: nginx
spec:
containers:
- name: nginx
image: nginx:latest
ports:
- containerPort: 80
 - Run below command to execute the file
kubectl apply -f "file_name"
 - To enter into the pod use command:
kubectl exec -it nginx-pod -n nginx -- bash
-

2. Namespaces

A Namespace is a way to divide cluster resources between multiple users or teams. They are virtual clusters backed by the same physical cluster.

Key Points:

- Namespaces are intended for environmental separation (e.g., dev, test, prod).
- Resources in one namespace cannot directly see or access resources in another namespace unless explicitly configured.
- Some resources like nodes, persistent volumes are not namespaced.
- By default, Kubernetes includes:
 - default: For objects without a specified namespace.
 - kube-system: For system components.
 - kube-public: Readable by all users.
- Namespaces help with:
 - Multi-tenancy
 - Resource Quotas

- Isolation and Access Control

Commands:

- If Minikube – *minikube start*
- *kubectl get namespaces*
- To get info about current serving cluster - *kubectl cluster-info*
- *kubectl create namespace my-namespace*
- Other way to create namespace via configuration file –

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: mysql-config
  namespace: my-namespace
data:
  db_url: mysqlservice.database
```

where “mysql-service” is database service url and “database” is the namespace name

- To list all resources bound to a namespace – *kubectl api-resources – namespaces=true*
 - To list all resources that are not bound to a namespace – *kubectl api-resources –namespaced=false*
 - Example – Volumes and nodes
 - Creating and running namespace configuration file - *kubectl apply -f mysqlconfig.yml*
 - *kubectl get configmap -n my-namespace*
 - For easy handling install kubectx on machine
 - *sudo apt install kubectx*
 - to check all the namespaces –
kubens //the default namespace will be highlighted
 - to set new active namespace –
kubens “new_namespace”
 - To check whether it is running or not:
kubectl port-forward pod/nginx-pod 8080:80 -n nginx
curl 127.0.0.1:8080
 - To get info about a particular pod:
kubectl describe pod/nginx-pod
-

3. Deployments

A Deployment is a higher-level object that manages Pods and ensures the desired number of Pods are always running.

Key Points:

- A Deployment:
 - Creates ReplicaSets.
 - ReplicaSets create and manage the Pods.
- Used to:
 - Roll out new versions of an application.
 - Roll back to previous versions if needed.
 - Scale applications up or down.
- Supports rolling updates (zero downtime deployments).
- Automatically replaces Pods if they fail.

Deployment Strategies:

- RollingUpdate (default): Gradually replaces old Pods with new ones.
- Recreate: Shuts down all old Pods before starting new ones.

Commands:

- Create a deployment.yml file for configuration

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: nginx-deployment
  namespace: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx

  template:
    metadata:
      name: nginx-dep-pod
```

labels:
app: nginx

spec:
containers:
- name: nginx
image: nginx:latest
port: 80

- *kubectl apply -f deployment.yml*
- *kubectl get deployment -n nginx* //Use only when accessing from another namespace
OR
- *kubectl get deployment*
- *kubectl get pods*
OR
- *kubectl get pods -n nginx* //Use only when accessing from another namespace
- To scale up the deployments:
kubectl scale deployment/nginx-deployment -n nginx --replicas=5 and then check the pods with the help of *kubectl get pods OR kubectl get pods -n nginx*
- Rolling updates commands:
kubectl set image deployment/nginx-deployment -n nginx nginx=nginx:1.27.3
 1. *kubectl* – CLI Kubernetes
 2. *set image* – command to update the image of running container
 3. *deployment/nginx-deployment* – targeting the nginx-deployment in the cluster
 4. *-n nginx* – specifies the namespace nginx
 5. *Nginx nginx=nginx:1.27.3* – updates the container nginx in the deployment to use the nginx version 1.27.3
 6. *kubectl get pods*

Summary Table

Concept	Purpose	Key Feature
Pod	Smallest deployable unit	Can have one or more containers
Namespace	Logical separation in cluster	Isolation, multi-tenancy
Deployment	Manages Pod lifecycle	Scaling, updates, rollback support

- *kubectl* context usually defaults to the default namespace.
- When creating resources, if namespace is not specified, it goes to the default namespace.

- You can switch namespaces using:
 - `kubectl config set-context --current --namespace=dev`
-

4. Replicasets:

1. What is a ReplicaSet?

A ReplicaSet (RS) is a Kubernetes controller that ensures a specified number of pod replicas are running at any given time.

Key Points:

- ReplicaSet maintains pod availability.
- If a pod crashes or is deleted, the ReplicaSet automatically creates a replacement pod.
- ReplicaSets are often not created directly by users. They are automatically created and managed by Deployments.
- Standalone ReplicaSets can be used, but Deployments are preferred as they offer more advanced features like rolling updates and rollbacks.

2. Why Use ReplicaSets?

- Ensures high availability by maintaining the desired number of pods.
- Provides self-healing: replaces failed pods automatically.
- Useful for scaling applications (adjust the number of replicas easily).

3. Core Concepts

Selector

- Used to match the ReplicaSet to the pods it should manage (based on labels).
- The ReplicaSet monitors only the pods with matching labels.

Labels

- Attached to pods.
- Identifies which pods belong to which ReplicaSet.

Template

- Describes how the new pods should be created (just like in a Deployment).
- Includes the pod specification: image, container, ports, etc.

4. ReplicaSet YML Example

```
kind: ReplicaSet
apiVersion: apps/v1
metadata:
  name: nginx-replicaset
  namespace: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx

  template:
    metadata:
      name: nginx-rep-pod
    labels:
      app: nginx

    spec:
      containers:
        - name: nginx
          image: nginx:latest
      ports:
        - containerPort: 80
```

Explanation:

- replicas: 3 → Desired number of pods.
- selector → Matches pods with label app: nginx.
- template → Describes the pod to be created.

5. Important Commands

Command	Purpose
<code>kubectl apply -f replicaset.yml</code>	Create a ReplicaSet
<code>kubectl get rs</code>	List all ReplicaSets
<code>kubectl describe rs nginx-replicaset</code>	Show detailed ReplicaSet info
<code>kubectl scale rs nginx-replicaset --replicas=5</code>	Scale ReplicaSet to 5 pods
<code>kubectl delete rs nginx-replicaset</code>	Delete the ReplicaSet

6. How ReplicaSets Work Internally

1. When you create a ReplicaSet, Kubernetes compares the desired replicas with the currently running pods.
2. If fewer pods are running → New pods are created.
3. If more pods are running → Extra pods are terminated.
4. ReplicaSets continuously monitor the state to maintain the desired number.

7. Difference Between ReplicaSet and Deployment

Feature	ReplicaSet	Deployment
Purpose	Maintain pod replicas	Manage ReplicaSets, support updates and rollbacks
Rolling Updates	Not supported	Supported
Rollbacks	Not supported	Supported
Manual Scaling	Supported	Supported
Preferred Use	Rarely used directly	Commonly used

Summary:

- Deployment → Full lifecycle manager (recommended)
- ReplicaSet → Just keeps pods alive and at desired count

8. Self-Healing Example

- Desired replicas: 3
- If one pod is deleted → ReplicaSet immediately creates a new pod to restore the count to 3.

- ReplicaSet constantly watches the pods with matching labels.

9. Selector and Label Relationship

- The selector in ReplicaSet is what tells it which pods to manage.
- Only pods with labels matching the selector will be controlled by the ReplicaSet.

Example:

selector:

matchLabels:

app: nginx

- All pods with the label app: nginx will be controlled by this ReplicaSet.

10. Summary Table

Concept	Purpose	Key Feature
Pod	Smallest deployable unit	Runs one or more containers
Namespace	Logical separation in cluster	Isolation, multi-tenancy
Deployment	Manages Pod lifecycle	Rolling updates, rollback
ReplicaSet	Maintains desired number of pods	Self-healing, scaling

Additional Notes:

- Do not directly update pods created by a ReplicaSet. The RS will detect a difference and create new pods to match the original specification.
 - ReplicaSets are almost always used as part of a Deployment in real-world Kubernetes setups.
-

5. Daemonsets:

1. What is a DaemonSet?

A DaemonSet is a Kubernetes controller that ensures a copy of a specific pod runs on every node (or on selected nodes) in the cluster.

Key Points:

- Ensures one pod per node.
- Automatically runs the pod on new nodes when they are added to the cluster.
- Commonly used for infrastructure-related tasks that need to run across all nodes.

2. Why Use DaemonSets?

DaemonSets are typically used to deploy:

- Log collection agents (e.g., Fluentd, Logstash)
- Monitoring agents (e.g., Prometheus Node Exporter)
- Networking components (e.g., Calico, Cilium)
- Storage daemons (e.g., GlusterFS, Ceph)
- Security agents (e.g., anti-virus scanners, intrusion detection)

3. How DaemonSets Work

- When a new node is added to the cluster → The DaemonSet automatically schedules a pod on it.
- When a node is removed → The pod is also removed.
- DaemonSets run one pod per node by default (unless node selectors or taints/tolerations are used to control where pods can run).

4. DaemonSet YAML Example

```
kind: DaemonSet
apiVersion: apps/v1
metadata:
  name: nginx-daemonsets
  namespace: nginx
spec:
  selector:
    matchLabels:
      app: nginx

  template:
    metadata:
      name: nginx-dmn-pod
    labels:
      app: nginx

    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Explanation:

- kind: DaemonSet → This resource is a DaemonSet.
- selector → Matches pods that the DaemonSet should manage.
- template → Describes the pod that should be created on each node.

5. Key Characteristics of DaemonSets

Feature	Description
Pod Scheduling	One pod per node automatically
New Node Behavior	Automatically schedules a pod on the new node
Node Removal Behavior	Pod is automatically removed
Manual Scaling	Not applicable (automatically scales per node)
Deletion Behavior	Deletes all DaemonSet pods when DaemonSet is deleted

6. Important DaemonSet Commands

Command	Purpose
<i>kubectl apply -f daemonset.yml</i>	Create the DaemonSet
<i>kubectl get daemonset</i>	List all DaemonSets
<i>kubectl describe daemonset nginx-daemonset</i>	Get detailed DaemonSet information
<i>kubectl delete daemonset nginx-daemonset</i>	Delete the DaemonSet and its pods
<i>kubectl get pods -o wide</i>	See on which nodes the pods are running

USES OF ALL SETS:

DAEMONSETS: At least one replica should be running on each and every node.

REPLICASETS: It will create the replicas of pods according to the number you give.

STATEFULSETS: It will keep the same states of all the replicas.

DEPLOYMENT: Rolling updates, scaling, auto scaling auto healing is done here.

6. JOBS

1. Kubernetes Job – Overview

A Job in Kubernetes ensures that a specified number of pods successfully complete a task. Unlike Deployments or DaemonSets that run continuously, a Job runs to completion.

Key Points:

- Runs pods to completion.
- Ensures the desired number of completions (successful runs).
- Useful for one-time tasks (e.g., data processing, backups, batch jobs).
- If a pod fails, the Job automatically starts a new pod to retry the task.
- Jobs can run single or parallel tasks.

Job YAML Example:

```
kind: Job
apiVersion: batch/v1
metadata:
  name: demo-job
  namespace: nginx
spec:
  completions: 1
  parallelism: 1
  template:
    metadata:
      name: demo-job-pod
    labels:
      app: batch-task
    spec:
      containers:
        - name: batch-container
          image: busybox:latest
          command: ["sh","-c","echo Jobs demo. && sleep 10"]
```

restartPolicy: Never

Job Key Fields:

Field	Purpose
completions	Number of successful runs required
backoffLimit	Number of retries on pod failure
restartPolicy	Should always be Never for Jobs
template	Pod specification

2. Job Behavior & Execution

Situation	Job Response
Pod completes successfully	Counts toward completions
Pod fails	Retries (up to backoffLimit)
All completions met	Job marked as successful
All retries exhausted	Job marked as failed

3. Important Job Commands

Command	Purpose
<i>kubectl apply -f job.yml</i>	Create the Job
<i>kubectl get jobs</i>	List all Jobs
<i>kubectl describe job example-job</i>	Detailed Job info
<i>kubectl delete job example-job</i>	Delete the Job
<i>kubectl get pods --selector=job-name=example-job</i>	List pods related to the Job
<i>kubectl logs pod/[pod_name]</i>	To check the logs of pods

4. Job Parallelism & Completion Modes

Jobs can run:

- Single Pod Jobs: Runs one pod to completion.
- Parallel Jobs:

- Run multiple pods simultaneously.
- Use parallelism and completions fields.

spec:

parallelism: 3

completions: 6

3 pods run simultaneously, job finishes after 6 completions.

5. Advanced Job Configurations

- RestartPolicy: Never → The pod will not be restarted by itself.
- ActiveDeadlineSeconds: Optional time limit for the job.
- TTLSecondsAfterFinished: Automatically cleans up job after completion.

6. When to Use a Job

- Batch Processing
 - Data Migrations
 - Database Backups
 - Image Processing
 - One-time Data Analysis
-

7. Kubernetes CronJob

A CronJob in Kubernetes is like a Linux cron—it runs Jobs on a scheduled, recurring basis.

Key Points:

- Runs Jobs on a predefined schedule.
- Schedule format uses cron syntax (minute, hour, day, etc.).
- Automatically creates a new Job at each scheduled time.
- Each scheduled Job runs independently.

CronJob YAML Example:

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: minute-backup
  namespace: nginx
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          name: minute-backup
        labels:
          app: minute-backup
        spec:
          containers:
            - name: backup-container
              image: busybox
              command:
                - sh
                - -c
```



```

- >
  echo "Backup Started" ;
  mkdir -p /backups &&
  mkdir -p /demo-data &&
  cp -r /demo-data /backups &&
  echo "Backup Completed" ;
volumeMounts:
  - name: data-volume
    mountPath: /demo-data
  - name: backup-volume
    mountPath: /backups
restartPolicy: OnFailure
volumes:
  - name: data-volume
    hostPath:
      path: /demo-data
      type: DirectoryOrCreate
  - name: backup-volume
    hostPath:
      path: /backups
      type: DirectoryOrCreate

```

CronJob Key Fields:

Field	Purpose
schedule	Cron format schedule (e.g. "0 * * * *")
jobTemplate	The job definition to be run
successfulJobsHistoryLimit	Retain X successful jobs (optional)
failedJobsHistoryLimit	Retain X failed jobs (optional)
concurrencyPolicy	Allow, Forbid, Replace overlapping jobs
startingDeadlineSeconds	Deadline for starting missed job runs

8. Cron Syntax Quick Reference

Cron Syntax Meaning

* * * * *	Every minute
* / 5 * * * *	Every 5 minutes
0 0 * * *	Every day at midnight
0 * * * *	Every hour

9. Important CronJob Commands

Command	Purpose
<i>kubectl apply -f cronjob.yml</i>	Create the CronJob
<i>kubectl get cronjobs</i>	List all CronJobs
<i>kubectl describe cronjob hello-cronjob</i>	Get detailed CronJob info
<i>kubectl delete cronjob hello-cronjob</i>	Delete the CronJob
<i>kubectl get jobs</i>	List Jobs created by the CronJob
<i>kubectl logs <job-pod-name></i>	View logs from the job pod

10. Concurrency Policies for CronJobs

Policy Behavior

Allow Allows concurrent runs of the job

Forbid Skips new job if the previous one is still running

Replace Cancels the currently running job and starts a new one

concurrencyPolicy: Forbid

11. Advanced CronJob Settings

- `startingDeadlineSeconds`: Time window in seconds to start a missed job if the CronJob controller was down.
- `successfulJobsHistoryLimit`: Limits how many successful jobs are kept in history.
- `failedJobsHistoryLimit`: Limits how many failed jobs are kept in history.
- `Suspend`: Temporarily disables the CronJob schedule.

12. Common CronJob Use Cases

- Daily Database Backups
- Periodic Log Rotation
- Regular Batch Data Imports
- Cleanup Scripts
- Scheduled Report Generation

13. Job vs CronJob vs Deployment vs DaemonSet Summary Table

Feature	Job	CronJob	Deployment	DaemonSet
Purpose	Run a task to completion	Run tasks on schedule	Long-running applications	One pod per node
Run Mode	Once	Recurring	Always running	Always running
Scheduling	Immediate	Cron-based	N/A	One per node
Scalability	Parallel supported	Jobs Creates automatically	Jobs Manual or auto scaling	One per node
Pod Restart Policy	Never	Never	Always	Always

14. Quick Memory Recap:

- Job: Run once, wait for successful completion.
 - CronJob: Run Jobs on a schedule.
 - Deployment: Manage long-running apps, scaling and rolling updates.
 - DaemonSet: Ensure a pod runs on every node.
-